

AI Interaction Log

Author: Bryce **Date:** 2026-09-25 **AI Tool:** Claude (Cowork)

Task / Prompt

Produce a paste-ready GitHub issue body for the authentication & core API endpoints task — the companion backend issue that sits on top of the database schema issue, in the same paste-ready format.

Context / Findings

- This issue depends on the database schema issue and is built on the same Sprint 1 tables (`users`, `profiles`, `resumes`).
- Endpoint scope was cross-checked against the Sprint 1 demo features (Registration/Login and Resume/Profile).
- Three decisions were left open for the team: (1) backend framework/language; (2) JWT vs. server-side sessions; (3) file upload handling (size limits, allowed types).

AI Output / Actions

- **Title:** [Task] Implement authentication & core API endpoints (Sprint 1)
- **Labels:** task, backend, api, sprint-1, priority: high
- **Depends on:** the database schema issue.
- **Endpoints — Auth:** POST `/auth/signup`, POST `/auth/login`, POST `/auth/logout`, GET `/auth/me`.
- **Endpoints — Profile:** GET `/profile`, PUT `/profile` (`full_name`, `headline`, `location`, `phone`, `bio`).
- **Endpoints — Resume:** POST `/resumes` (upload + type/size validation), GET `/resumes`, DELETE `/resumes/:id`.
- **Tasks:** choose framework/language + document; project skeleton + env config + DB connection; password hashing (bcrypt/argon2); auth endpoints w/ token sessions; profile endpoints; resume endpoints w/ validation; auth middleware on protected routes; input validation + consistent error responses; basic auth-flow tests; endpoint documentation.
- **Definition of Ready:** DB schema merged/available; framework agreed; auth strategy decided (JWT vs. session); resume storage approach decided.
- **Definition of Done:** endpoints implemented + testable; passwords hashed + protected routes enforce auth; upload validates type/size; auth-flow tests pass; endpoints documented; merged via reviewed PR.

Decision

Bryce: Used the issue body mostly as-is, with a few wording tweaks for clarity. Kept the three open decisions (framework, JWT vs sessions, upload handling) flagged rather than deciding them myself, since those are calls for the whole team to make once we're all in the same room. A couple of the tasks might get trimmed depending on how far we actually get this sprint.

Reflection

Bryce: I don't have much backend/API experience, so having the endpoints and the definition of ready/done laid out gave me a clear picture of what building the auth layer actually involves. It also made the dependency on the database issue obvious, which helped me sequence the work for whoever picks it up.